

Responding to the global challenge of climate change.

Learn more

MIT
Technology Review

Featured Topics Newsletters Events Podcasts

SIGN IN

SUBSCRIBE

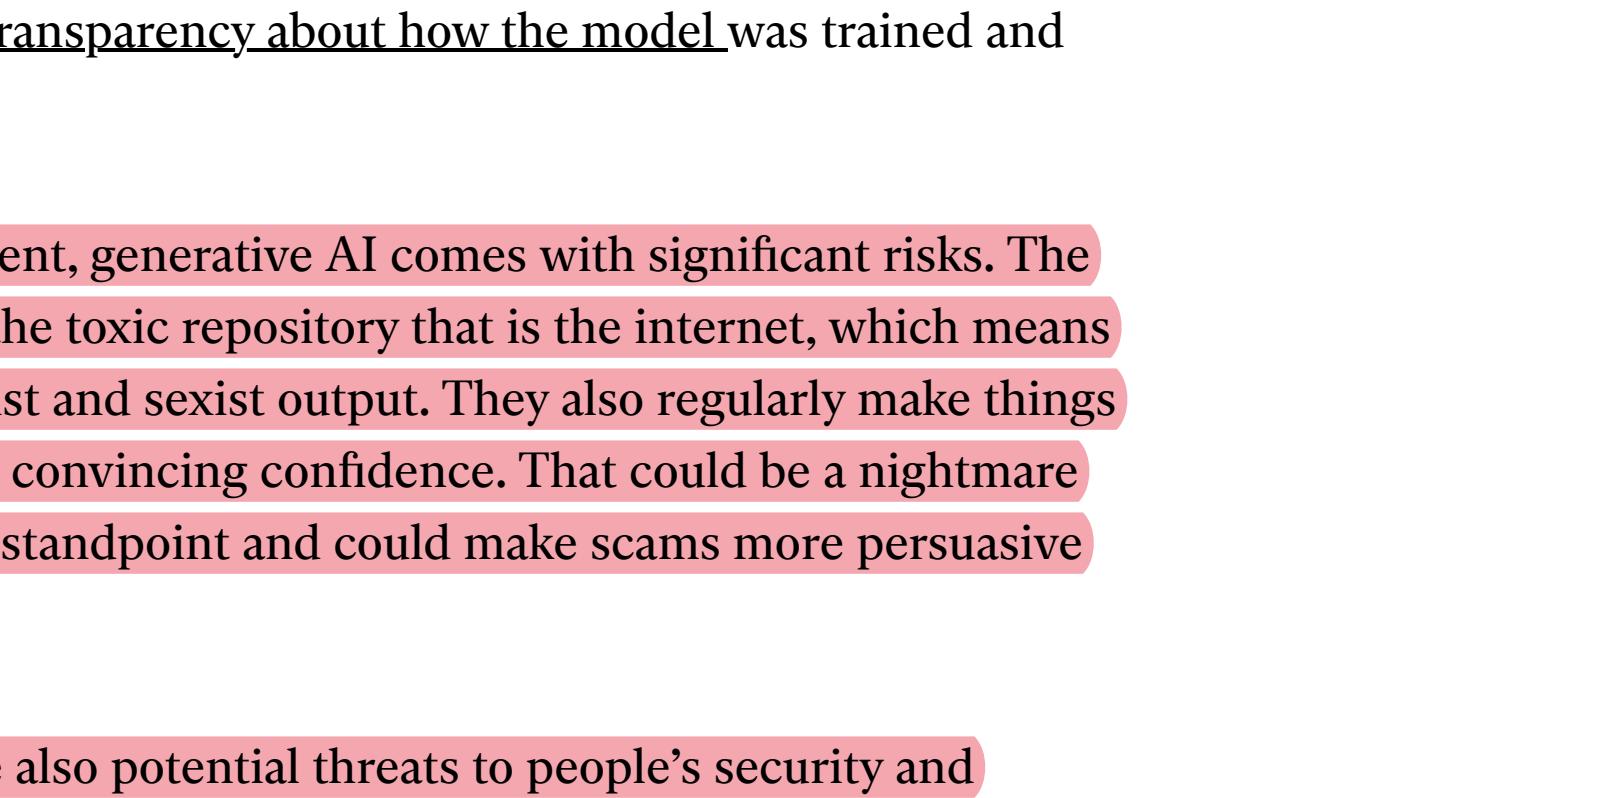
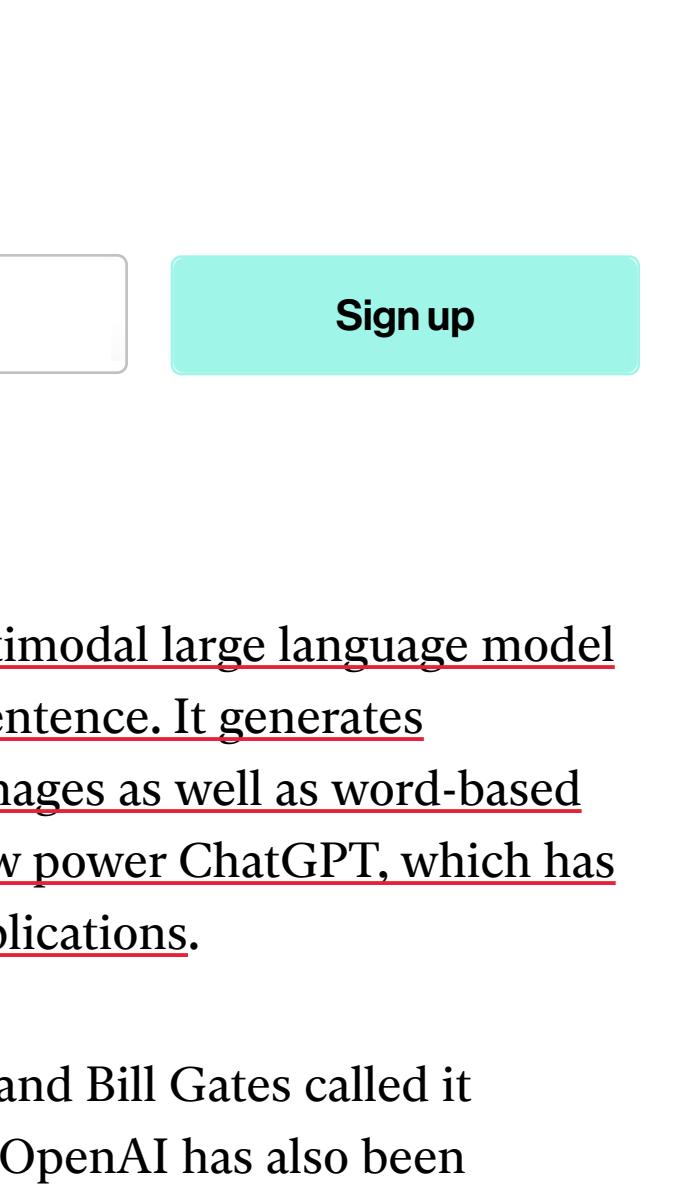
TECH POLICY

An early guide to policymaking on generative AI

How lawmakers are thinking about the risks of the latest tech revolution

By Tate Ryan-Mosley

March 27, 2023



ARI LILIAN

This article is from *The Technocrat*, MIT Technology Review's weekly tech policy newsletter about power, politics, and Silicon Valley. To receive it in your inbox every Friday, sign up here.

Earlier this week, I was chatting with a policy professor in Washington, DC, who told me that students and colleagues alike are asking about GPT-4 and generative AI: What should they be reading? How much attention should they be paying?

She wanted to know if I had any suggestions, and asked what I thought all the new advances meant for lawmakers. I've spent a few days thinking, reading, and chatting with the experts about this, and my answer morphed into this newsletter. So here goes!

Though GPT-4 is the standard bearer, it's just one of many high-profile generative AI releases in the past few months: Google, Nvidia, Adobe, and Baidu have all announced their own projects. In short, generative AI is the thing that everyone is talking about. And though the tech is not new, its policy implications are months if not years from being understood.

WEEKEND READS

Technology in perspective.

Enter your email

Sign up

By signing up, you agree to our [Privacy Policy](#).

GPT-4, released by OpenAI last week, is a multimodal large language model that uses deep learning to predict words in a sentence. It generates remarkably fluent text, and it can respond to images as well as word-based prompts. For paying customers, GPT-4 will now power ChatGPT, which has already been incorporated into commercial applications.

The newest iteration has made a major splash, and Bill Gates called it "revolutionary" in a letter this week. However, OpenAI has also been criticized for a lack of transparency about how the model was trained and evaluated for bias.

Despite all the excitement, generative AI comes with significant risks. The models are trained on a toxic repository that is the internet, which means they often produce racist and sexist output. They also regularly make things up and state them with convincing confidence. That could be a nightmare from a misinformation standpoint and could make scams more persuasive and prolific.

Generative AI tools are also potential threats to people's security and privacy, and they have little regard for copyright laws. Companies using generative AI that has stolen the work of others are already being sued.

Alex Engler, a fellow in governance studies at the Brookings Institution, has considered how policymakers should be thinking about this and sees two main types of risks: harms from malicious use and harms from commercial use. Malicious uses of the technology, like disinformation, automated hate speech, and scamming, "have a lot in common with content moderation," Engler said in an email to me, "and the best way to tackle these risks is likely platform governance." (If you want to learn more about this, I'd recommend listening to this week's *Sunday Show* from Tech Policy Press, where Justin Hendrix, an editor and a lecturer on tech, media, and democracy, talks with a panel of experts about whether generative AI systems should be regulated similarly to search and recommendation algorithms. Hint: Section 230.)

Policy discussions about generative AI have so far focused on that second category: risks from commercial use of the technology, like coding or advertising. So far, the US government has taken small but notable actions, primarily through the Federal Trade Commission (FTC). The FTC issued a warning statement to companies last month urging them not to make claims about technical capabilities that they can't substantiate, such as overstating what AI can do. This week, on its business blog, it used even stronger language about risks companies should consider when using generative AI.

"If you develop or offer a synthetic media or generative AI product, consider at the design stage and thereafter the reasonably foreseeable—and often obvious—ways it could be misused for fraud or cause other harm. Then ask yourself whether such risks are high enough that you shouldn't offer the product at all," the blog post reads.

The US Copyright Office also launched a new initiative intended to deal with the thorny policy questions around AI, attribution, and intellectual property.

The EU, meanwhile, is sticking true to its reputation as the world leader in tech policy. At the start of this year my colleague Melissa Heikkila wrote about the EU's efforts to try to pass the AI Act. It's a set of rules that would prevent companies from releasing models into the wild without disclosing their inner workings, which is precisely what some critics are accusing OpenAI of with the GPT-4 release.

The EU intends to separate high-risk uses of AI, like hiring, legal, or financial applications, from lower-risk uses like video games and spam filters, and require more transparency around the more sensitive uses. OpenAI has acknowledged some of the concerns about the speed of adoption. In fact, its own CEO, Sam Altman, told ABC News he shares many of the same fears. However, the company is still not disclosing key data about GPT-4.

The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

Hyper-realistic beauty filters are here to stay

A new filter on TikTok has the internet up in arms. It's an important debate for anyone who cares about the future of social media.

By Tate Ryan-Mosley

AI might not steal your job, but it could change it

AI is already being used in the legal field. Is it really ready to be a lawyer?

By Tate Ryan-Mosley

Illustration by Rose Wong

Get the latest updates from MIT Technology Review

Discover special offers, top stories, upcoming events, and more.

Enter your email



Privacy Policy

Cookie settings



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

DEEP DIVE

TECH POLICY

The internet is about to get a lot safer

Europe's big tech bill is coming to fruition. Here's what you need to know.

By Tate Ryan-Mosley

Cookie Policy

We use cookies to give you a more personalized browsing experience and analyze site traffic. See our [cookie policy](#).

Accept all cookies

Cookies settings

View offers

Sign in

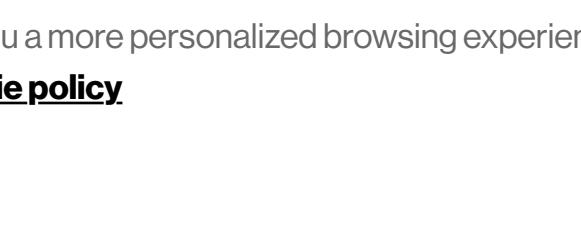


The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

STAY CONNECTED



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

DEEP DIVE

TECH POLICY

The internet is about to get a lot safer

Europe's big tech bill is coming to fruition. Here's what you need to know.

By Tate Ryan-Mosley

Cookie Policy

We use cookies to give you a more personalized browsing experience and analyze site traffic. See our [cookie policy](#).

Accept all cookies

Cookies settings

View offers

Sign in

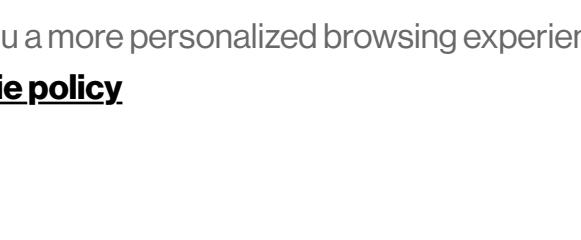


The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

STAY CONNECTED



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

DEEP DIVE

TECH POLICY

The internet is about to get a lot safer

Europe's big tech bill is coming to fruition. Here's what you need to know.

By Tate Ryan-Mosley

Cookie Policy

We use cookies to give you a more personalized browsing experience and analyze site traffic. See our [cookie policy](#).

Accept all cookies

Cookies settings

View offers

Sign in

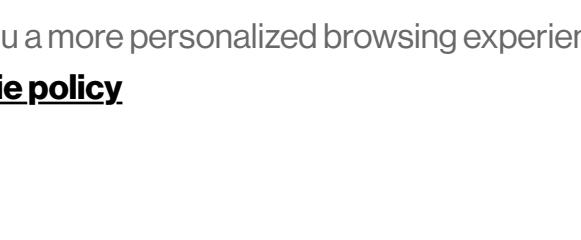


The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

STAY CONNECTED



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

DEEP DIVE

TECH POLICY

The internet is about to get a lot safer

Europe's big tech bill is coming to fruition. Here's what you need to know.

By Tate Ryan-Mosley

Cookie Policy

We use cookies to give you a more personalized browsing experience and analyze site traffic. See our [cookie policy](#).

Accept all cookies

Cookies settings

View offers

Sign in

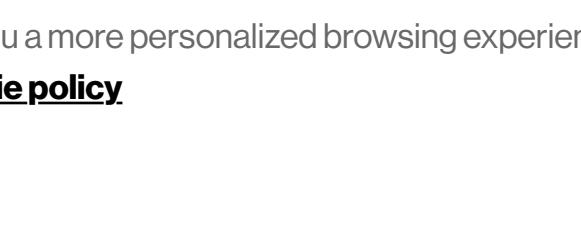


The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

STAY CONNECTED



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

DEEP DIVE

TECH POLICY

The internet is about to get a lot safer

Europe's big tech bill is coming to fruition. Here's what you need to know.

By Tate Ryan-Mosley

Cookie Policy

We use cookies to give you a more personalized browsing experience and analyze site traffic. See our [cookie policy](#).

Accept all cookies

Cookies settings

View offers

Sign in

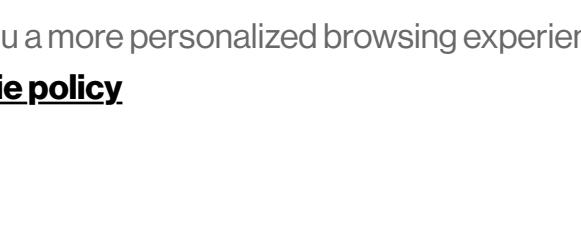


The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

STAY CONNECTED



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

DEEP DIVE

TECH POLICY

The internet is about to get a lot safer

Europe's big tech bill is coming to fruition. Here's what you need to know.

By Tate Ryan-Mosley

Cookie Policy

We use cookies to give you a more personalized browsing experience and analyze site traffic. See our [cookie policy](#).

Accept all cookies

Cookies settings

View offers

Sign in



The Design Issue

Subscribe now to explore our latest issue and learn how design influences AI and other technologies.

See offers

STAY CONNECTED